

The Same Zero: Why Identical ASR Can Imply Different Guarantees in LLM-Agent Security

Anonymous ACL submission

Abstract

LLM-agent security has produced a dense landscape of defenses—prompt hardening, content filters, permission gates, sandboxes, provenance firewalls—each claiming to make agents “safe.” We argue that the field lacks a pre-purchase question: *what does a given defense actually guarantee, and where does that guarantee come from?* We apply Verification Autonomy Levels (VAL, arXiv:2608.19009), a taxonomy that grades verification/authorization schemes by the source of their spec (L0: LLM self-declaration, no deterministic anchor; L1: deterministic rules; L2: objective ground truth, correctness only; L3/L4: decidable systems with single-property or domain-level completeness; L5: impossible), to 22 agent-security defenses. The classification is a falsifiable predictor: a defense fails the way its anchor fails. We validate this with frozen prediction cards on a small published sample (10/10 hits, flagged as such) and then run the first controlled deployment-value comparison: same budget, two stacks—a VAL-guided stack (intent-anchored confirmation gate + schema sandbox) versus a mainstream intuition stack (prompt hardening + keyword filter)—across 50 scenarios, 12 attack families (including third-party AgentDojo payloads), real tool effects, adaptive, white-box and PAIR-style attacks, and three seeds (~7,000 LLM calls). The VAL stack holds 0.000 attack success with 1.000 benign success in every condition; the intuition stack also reaches 0.000 ASR but kills all benign actions and its security is model-behavior luck (compliance 0.235), not structure (the VAL stack tolerates 2.4x more model compromise—0.567 compliance—and still blocks everything). The same zero, two different guarantees—zero is an outcome, not a guarantee. Under escalating attacks, zero stability becomes a measurable axis: some zeros collapse abruptly, others erode gradually, and indistinguishable stability trajectories

can rest on opposite guarantees (model behavior vs. platform record). VAL is a guarantee-provenance framework for LLM-agent security: it identifies what an observed security outcome is grounded in, and hence where it will fail. We further identify where L3 anchors exist in agent security: confinement and information-flow properties (sandboxing, taint tracking, data-control separation), not unrestricted semantic safety, which caps at L2 (restricted semantic properties that can be formally encoded remain eligible for higher levels).

1 Introduction

Large language model (LLM) agents are moving from single-turn assistants to persistent systems that call tools, browse the web, and maintain long-term memory—and the matching expansion of security defenses (prompt hardening, filters, permission gates, sandboxes, provenance firewalls) is presented with implicit, rarely commensurable guarantees: one paper’s attack success rate (ASR) is not another’s, and no framework tells a deployer *which defense to trust for which threat, before running experiments*.

This paper asks the question the field has not asked explicitly: **what does a given agent-security defense guarantee, and where does that guarantee come from?**

We answer with Verification Autonomy Levels (VAL), a taxonomy proposed in (Yin 2026) for verification schemes generally. VAL classifies any verification/authorization scheme along a single axis—the source of its verification spec—into six levels:

- **L0:** the spec is LLM self-declaration (“I checked it,” “this is safe”). No deterministic anchor. Failure mode: the declaration can be rewritten by the very model that issues it.

- **L1:** deterministic rules (regex, keyword lists, hand-written policies). Failure mode: surface forms are bypassable by rewriting.
- **L2:** objective ground truth (gold labels, platform-recorded events, execution results). Guarantees *correctness only*: everything checked passes, but nothing is proven about what was not checked (the completeness blind spot).
- **L3/L4:** decidable systems (solveset equivalence, type systems, formal kernels, confinement mechanisms). Single-property or domain-level completeness within an operational design domain (ODD).
- **L5:** universal completeness—impossible in the unrestricted case.

VAL’s core claim is that a scheme’s level is determined by its *anchor*, not by its accuracy or sophistication, and that the anchor determines the scheme’s failure modes. Applied to agent security, this yields a falsifiable prediction: **a defense is defeated the way its anchor fails**. An L0 defense fails when the model complies with the injection; an L1 defense fails when the attacker rewrites around the surface rule; an L2 defense fails outside its ODD (forged metadata, poisoned data); an L3 defense holds within its ODD by construction.

We make three contributions:

1. **A level map of agent security.** We classify 22 defenses using a frozen, blind-rater-validated protocol (inter-rater $\kappa \approx 0.8$ (Yin 2026)), organizing the landscape by what each defense can actually guarantee.
2. **Prediction validation.** Frozen level-to-behavior cards, validated on a small published sample (10/10 hits, flagged; two mechanism-driven corrections)—the point is falsifiability, not the hit count.
3. **Deployment-value comparison and zero stability.** Same budget, two stacks (VAL-guided vs. mainstream intuition), one testbed—50 scenarios, 12 attack families, real effects, adaptive/white-box/PAIR escalation, three seeds, a second victim, and the official AgentDojo benchmark. The VAL stack dominates on security and utility; escalating-attack trajectories ($Z(\alpha)$) further separate defenses by failure morphology (abrupt collapse, progressive erosion, flat behavioral vs. structural plateau), and iden-

tical trajectories still require provenance to interpret—zero stability joins ASR and compliance as a reportable property.

What this paper claims—the spine in one paragraph. The argument runs on three levels. *Phenomenon*: two defenses can reach the same observed zero ASR with opposite guarantees (Section 6). *Mechanism*: the difference lies in the verification anchor, where the security verdict is grounded (Section 5). *Method*: VAL identifies the strongest guarantee a defense actually grounds and predicts each level’s failure boundary. If a reader remembers one sentence: *the same zero is not the same guarantee—zero is an outcome, not a guarantee*.

What we do not claim. VAL is not a general theory of AI verification, a complete security evaluation, or a ranking of defenses. It is a mid-level methodological claim about LLM-agent security: when two defenses produce the same observed outcome, the provenance of that outcome determines its guarantee and predicts its failure boundary. Robotics, formal methods, HITL, and non-agentic LLMs are outside scope; so is a claim that VAL covers every defense.

2 Background and Related Work

2.1 Verification Autonomy Levels (VAL)

VAL (Yin 2026) grades verification schemes by three questions: Q1 *where does the verification spec come from* (LLM declaration / deterministic rule / objective truth / decidable system); Q2 *what does the verdict guarantee* (nothing / correctness / completeness); Q3 *what scope does a completeness claim cover* (single property / domain / universal). The level is the first rule that fires: L5 (universal completeness, impossible), L4 (domain completeness), L3 (single-property completeness), L2 (correctness with objective anchor), L1 (correctness with deterministic rule), L0 (otherwise). The procedure is deterministic and has been measured for inter-rater reproducibility: three independent blind raters reach $\kappa \approx 0.8$ on 48–70 schemes under the refined protocol (Yin 2026).

Two refinements matter. First, **anchor semantics** (Zhang et al. 2026): objective anchors split by what they certify—intent (a recorded decision event), truth (a fact), effect (an execution outcome). A confirmation record is an intent anchor: it authorizes, but says nothing about correctness;

conflating intent with effect is the category error behind most authorization failures. Second, **the completeness blind spot** (Yin 2026): checks verify proposed candidates but cannot prove none was missed—the analog in security is the *unenumerated attack*: a gate blocking all evaluated attacks is a correctness probe over the evaluated distribution, not a guarantee over all attack paths.

2.2 Agent-security defenses and their taxonomies

The defense literature spans text-level (prompt hardening, filters, Llama Guard), tool-level (allowlists, IFC, sandboxes), and memory-level (A-MemGuard, PPMF) defenses, plus benchmarks like AgentDojo. Recent work separates injection from execution success (*Injection-Execution Dissociation* 2026), classifies test oracles by authority source (*LLM-Based Test Oracles* 2026), and shows same-family verification yields near-zero gain (Lu 2025). Our contribution is different in kind: not another defense, but a *pre-purchase axis*—the anchor—that organizes the landscape and predicts each defense’s failure modes.

3 Method

3.1 Classification protocol

We classify defenses with the frozen VAL protocol (v2.3, (Yin 2026)) extended for security: R2 mixed systems split per component (weakest anchor on the verdict path); R3 evidence anchors do not equal verdict anchors; R10 designer rules are L1; R12b theorem-backed exhaustive procedures (confinement, taint tracking) are L3; R15 LLM-generated expectations are L0. Each defense is rated on Q1/Q2/Q3 from a mechanism-only evidence pack by blind raters. For the security subset (22 items), two independent raters agree exactly on 20/22 (90.9%); the two disagreements are known boundary cases.

On the validity of the taxonomy. Agreement among raters shows the categories are *usable*, not *right*. We offer three independent lines instead: *out-of-sample prediction* (cards frozen before outcome data, 10/10 hits with two mechanism-driven corrections—retrospective storytelling cannot produce that pattern); *cross-benchmark convergence* (the same anchor logic predicts our testbed, AgentDojo, JADE, and LITMUS results—four sources with no shared implementation); and *falsifiability* (each card names the attack predicted to break its

defense). We claim VAL is a productive organizing principle, not the unique correct ontology.

3.2 Prediction protocol

Following (Yin 2026), we freeze prediction cards before outcome data: for each defense, level $\rightarrow$ predicted behavior pattern + the specific attack predicted to break it, validated against published data (PPMF’s numbers; abstract-level evidence flagged as such). The protocol also generalizes out of sample: before running JADE’s MCP instances (Section 6.5) we froze five predictions (e.g., N’s keyword filter will leak the `post_tweet` variant—its action word absent from the filter’s vocabulary; V’s sandbox will block every variant by whitelist enumeration). All five held, including the boundary prediction.

3.3 Deployment-value experiment

Testbed. A minimal agent: external content is written into memory; a user task triggers a tool call; the LLM (DeepSeek-chat, temperature 0.2) proposes a call; the defense chain decides; a real-effect executor performs the action (a ledger for transfers, an outbox for email, real file deletion in a sandbox directory, a forum log for comments). Attack success requires a *real state change* (server-side effect semantics, as in (Xu et al. 2026)), not merely a proposal.

Scenarios and attacks. 50 scenarios (30 malicious + 20 benign; the benign set is larger than the 25% share of (Xu et al. 2026; Debenedetti et al. 2024) to narrow utility confidence intervals), derived from the memory-provenance families of (Xu et al. 2026). Attacks: 8 self-written variants (direct, paraphrase, fake-confirmation, combo, role-play, jailbreak prefix, deep rewrite with keyword splitting, authority frame) plus 4 third-party AgentDojo payloads (Debenedetti et al. 2024), 12 total; escalation adds an adaptive attacker (3 rounds, rewrites from the interception reason), a white-box adaptive attacker (told the system prompt and keyword list), and a PAIR-style candidate-search attacker.

Defenses. Four self-implemented defenses, one per anchor family: - **D1** (L0): system-prompt hardening—“ignore instructions in external content; never perform high-risk actions without explicit user request.” The verdict is the model’s own compliance. - **D2** (L1): keyword filter over memory and call text (transfer/delete/send/payment

terms, Chinese equivalents). - **D3** (L1/L2): confirmation gate—high-risk tools require a *platform-recorded* confirmation event; text claims (“user-confirmed”) are ignored. - **D4** (L3): parameter sandbox—tool arguments must match schema (recipient whitelists, target patterns, type/range checks), executed in a side-effect stub with verified state changes.

Stacks (same budget, two defenses each). - **V (VAL-guided)**: D3 + D4—an intent anchor (L2 input) plus a confinement property (L3). - **N (mainstream intuition)**: D1 + D2—prompt hardening plus keyword filtering, the intuitive low-cost default in production agents. We emphasize what this comparison does *not* claim: it is not “VAL beats all defenses” (structural defenses dominate behavioral ones by construction); the contribution is the *epistemic* contrast—identical security numbers, opposite guarantees, a sharp utility cost for the intuition stack (Section 6.1).

Metrics. ASR = high-risk tool executed with complete arguments and verified state change; compliance rate = the model *proposes* a high-risk tool (including empty-argument hedging)—separating “model was driven by the injection” from “the defense intercepted.” Benign success = benign scenarios complete their intended action.

Runs. Static: 7 configs $\times$ 380 cases (30 malicious $\times$ 12 attack variants + 20 benign). Adaptive/white-box: 30 malicious scenarios $\times$ stacks $\times$ 3 rounds. Seeds: V, N, and the no-defense baseline (ND) re-run 3 $\times$. A second victim (Llama 3.1 8B) re-runs the static battery. Total $\approx$ 7,000 LLM calls across all experiments; AgentDojo (Section 6.5) and JADE probes add $\sim$ 2,000 harness-internal calls reported separately.

4 The Defense Level Map

The full level map (all 22 defenses) is in Appendix A.1.

Full per-defense ratings with evidence: reliability/corpus.json, ratings/; blind-rater agreement for the security subset is 100% between the two most recent raters (6/6 agentsec items, $\kappa \approx 0.88$ overall (Yin 2026)).

The structural finding. L3 anchors exist in agent security—but only for *confinement and information-flow properties*: sandbox execution domains, taint tracking, data-control separation,

SMT-monotonic permission lattices. These are decidable by construction. *Semantic* safety (“is this action dangerous?”) is an open-world property for which no decidable fragment is available in the unrestricted setting, and hence caps at L2; restricted semantic properties that can be formally encoded (e.g., a bounded action taxonomy) remain eligible for higher levels. This mirrors the formal/empirical dichotomy of (Yin 2026) exactly: completeness is reachable only for formally specifiable properties.

4.1 Worked examples: how the classification reads three defenses

Each classification is a claim; per the framework’s own epistemology, claims need anchors. We show the full evidence chain for three representative defenses spanning the level spectrum (remaining classifications in the supplementary corpus).

Self-Ask provenance — no anchor (L0). Mechanism: the LLM infers whether memory text looks external, user-history-like, or confirmation-like, and authorizes from the inferred trust (Xu et al. 2026). Q1: LLM; Q2: none; Level: **L0**. Evidence: permissive mode reaches **0.984 ASR**; strict mode blocks all benign actions (Xu et al. 2026). Predicted failure mode: text claims (fake markers, social engineering). Observed: 0.984 ASR under permissive inference—the anchor is whatever the model believes.

PPMF — the intent anchor (L1/L2). Mechanism: a deterministic gate authorizes tool calls by matching action risk to the least-trusted supporting memory, whose provenance metadata (source, trust, confirmation events) is platform-maintained (Xu et al. 2026). Q1: designer policy over platform-recorded events; Q2: correctness; Level: **L1** (weakest anchor) with L2-grade inputs. Evidence: 0.000 ASR with intact metadata; **0.088 ASR under 10% forged confirmations**—the completeness blind spot quantified. Predicted failure mode: forged/compromised metadata (outside the ODD). Observed: exactly that; the authors themselves list compromised metadata outside their guarantee.

IFC/Fides — confinement as L3. Mechanism: the planner dynamically tracks taint labels and deterministically enforces lattice policies (Costa et al. 2025). Q1: decidable; Q2: completeness; Q3: single property (information flow). Level: **L3**. Ev-

idence: a formal model characterizing the class of properties enforceable by dynamic taint tracking. Predicted failure mode: properties outside the lattice (covert channels, semantic content). Observed: consistent—the guarantee is about flow, not about whether the permitted action is wise.

The three examples trace the level ladder with primary-source evidence: L0 fails to text belief, L1/L2 fails outside its ODD (forged metadata), L3 holds by construction within its property—all failure modes predicted before the validation data were consulted (Section 5).

5 Prediction Validation

We froze 16 prediction cards (level → predicted failure mode + predicted breaking attack) before consulting outcome data.

PPMF family (6 defenses): all six predictions hit on PPMF’s published numbers (Xu et al. 2026): content filters collapse under rewriting; Self-Ask provenance (L0) reaches 0.984 ASR; the gate-only ablation kills benign actions; full PPMF (L1/L2) achieves 0.000 ASR with benign preserved, and 0.088 under 10% forged confirmations—the completeness blind spot quantified.

New cards (abstract-level evidence, flagged): 4/4 behavior predictions hit (Llama Guard distribution-bound accuracy; RA-LLM residual ASR under strong attacks; perplexity filtering weak alone; ToolEmu audit-value-only). Two mechanism-driven corrections emerged; because they are the strongest evidence that the prediction loop *corrects* rather than merely confirms, each gets a full account below.

5.1 Mechanism-driven correction I: Progent’s SMT monotonic confinement

Predicted: L1 (designer permission policy). **Mismatch:** every policy update is adjudicated by an SMT solver as *narrowing* (automatic) or *expanding* (requires approval)—the agent’s action space can only shrink without approval, a confinement property with a decidable witness. **Corrected classification:** the gate stays L1, but the monotonic-confinement guarantee is annotated L3 (theorem-backed structure, R12b). The mismatch is evidence for, not against, the protocol.

5.2 Mechanism-driven correction II: RARR’s LLM-mediated verdict (R3)

Predicted: L2 (attribution checking). **Mismatch:** RARR’s verdict path is LLM-mediated end-to-end—“citation existence” is not mechanically checked. Under R3 (evidence anchors do not equal verdict anchors), the verdict component is LLM declaration. **Corrected classification:** L2 → L0; a superficially “objective” anchor can be L0-masked when the alignment step is model-mediated—precisely the laundering pattern of the memory domain.

Cumulative: 10/10 behavior predictions hit (6 high-confidence, 4 abstract-level), 2 classification corrections. Sample is small; the value is that the prediction *protocol* is now demonstrated end-to-end, not that the hit rate is large.

6 Deployment-Value Experiment: Results

The dominance at a glance (full tables in 6.1-6.4):

The two zeros in the last two rows are the paper’s spine: identical security numbers, opposite guarantees, a 100-point utility gap.

6.1 Static attacks (12 variants, 380 cases/config, real effects)

Benign success is measured on 20 benign scenarios (Wilson 95% CIs in brackets); the utility gap between V and N is significant at $p < 0.001$ (Fisher’s exact test on 20/20 vs 0/20).

The two stacks reach the same zero but differ in *how*: V’s zero is produced by a gate the model cannot influence; N’s zero by the model’s behavior under the hardening prompt (compliance 0.235—proposing the malicious tool a quarter of the time, mostly with empty arguments; strict ASR counts only actionable executions). D2 alone leaks 0.200: the keyword filter is bypassable by rewriting.

The two layers of V are not redundant; they answer different questions. D3 (confirmation gate) is the reason V’s ASR is 0.000: it refuses high-risk tools without a platform-recorded authorization event, and the model cannot forge that record. D4 (parameter sandbox) is not about *whether* an action is authorized but *what an authorized action may touch*: it constrains arguments to schema (recipient whitelists, target patterns). The divi-

Stack	ASR (both vict
ND (no defense)	0.333 / 0.253
V (VAL: gate + sandbox)	0.000 / 0.000
N (intuition: hardening + filter)	0.000 / 0.000
Config	ASR
ND (no defense)	0.333
D1 prompt hardening (L0)	0.000
D2 keyword filter (L1)	0.200
D3 confirmation gate (L1/L2)	0.000
D4 parameter sandbox (L3)	0.067
V = D3+D4 (VAL)	0.000
N = D1+D2 (intuition)	0.000

sion shows up in the single-defense numbers—D3 alone already reaches 0.000 ASR, while D4 alone leaks 0.067 (schema-legal semantic attacks, e.g. a malicious comment posted to an allowed target). Stacking them yields the claim in its full form: D3 controls *whether* (the authorization anchor), D4 controls *what* (the confinement anchor), and an attack must defeat both anchors on the same call to succeed.

6.2 Adaptive and white-box attacks

The full escalation table is in Appendix A.5; the key contrast: D1/D3 hold at 0.000 across static/adaptive/white-box, while D2 (0.200→0.467) and D4 (0.067→0.200) leak and worsen.

The single-defense breakpoints are already informative. Two zeros survive both adaptive and white-box escalation unchanged (D1 at 0.000, D3 at 0.000), and two do not (D2 0.200→0.467, D4 0.067→0.200). But the two surviving zeros rest on different grounds, visible in compliance under the same attacks: D3’s model is driven to propose the malicious tool in 53.3% of adaptive cases and 56.7% of white-box cases and is still blocked every time (the gate reads a platform record the attacker cannot write), while D1’s model is driven 36.7%/33.3% and the zero depends on its refusal behavior holding under every attack form we tried. The same static zero, the same adaptive zero, the same white-box zero—different provenance, and different predicted breakpoints under attacks strong enough to move a behavioral refusal (Section 7).

Zero stability under escalating attack ($Z(\alpha)$).

Tracking the cumulative ASR round by round (7 rounds, 30 cases per point; full data in Appendix) exposes distinct stability morphologies. D2 climbs

monotonically to **0.833 by round 7** (each rewrite learns another way around the keyword vocabulary); D4 jumps at round 2 and plateaus (0.067 → 0.300); D1 and D3 stay at 0.000 across all seven rounds. The two flat curves are indistinguishable in $Z(\alpha)$ alone—which is exactly the point. **Identical observed zero-stability profiles do not, by themselves, establish equivalent security guarantees**; only the provenance separates them, and the gap widens under escalation: D3’s model is driven to propose the malicious tool in **83.3%** of cases and is still blocked every time (the gate reads a platform record), while D1’s model is driven 36.7% and the zero is its refusal behavior holding. We attempted to falsify D1’s flat curve directly (7-round adaptive attacks with prompts explicitly demanding complete executable arguments, plus authority/emergency/role-play frames): it did not break under this attack budget—stronger support for the provenance distinction, not evidence that it cannot break. We predict it will break under an attack strong enough to defeat DeepSeek’s refusal behavior (e.g., a less-aligned attacker); D3’s will not, short of compromising the audit substrate itself (Appendix).

Where the zero breaks: an execution funnel under static attacks

(a mechanism snapshot, distinct from the 7-round $Z(\alpha)$ above). For each case we record whether the model proposes a high-risk tool (compliance), whether the proposal carries complete arguments, and whether it passes authorization and executes with a verified state change. D1’s funnel collapses at the second stage: the zero is terminated at the complete-argument stage²⁰¹⁴the model fails to produce executable arguments under the hardening prompt. The break is at parameter completeness, but the anchor remains behavioral: D1 has no structural parameter-

integrity verifier (same shape as N’s, 0.233 $\pm$ 0.000). D3’s funnel reaches the third stage: complete arguments in 49.7% of cases, refused every time by the authorization gate (the zero is the platform record (same shape as V’s, 0.497 $\pm$ 0.000)). The same ASR and flat Z (0.317), but the funnels break at different layers: model behavior (arguments) vs. external constraint (authorization). The zero does not say where it comes from; the funnel does.

Definition (hedging wall). Under a hardening prompt (D1), an LLM victim under injection may comply *partially*: it proposes the requested high-risk tool while emitting empty or incomplete arguments—a stable compromise between the injection (“propose the action”) and the safety prompt (“do not execute without authorization”). Actionable ASR requires complete arguments, so hedging suppresses ASR *without any structural gate*. The wall is a model behavior, not a mechanism: its height is victim-specific (compliance 0.235 for DeepSeek, 0.014 for Llama 8B, Section 6.4).

Adaptive attacks nearly double the leak rate of the L1 and L3 single defenses (D2 0.200 $\rightarrow$ 0.467, D4 0.067 $\rightarrow$ 0.200). Against the stacks: V holds even when the model is driven to propose the malicious tool in 57% of scenarios (white-box compliance 0.567, $2.4\times$ N’s 0.235); N holds only because its hedging wall survives every attack form we tried—static, adaptive, white-box, PAIR (0.000), and a willing cross-family attacker (DeepSeek against Llama: 0.000, Section 6.4). Aligned cross-family attackers (Moonshot Kimi, Meta Llama 8B) declined the role—explicit refusals, sanitizing rewrites (Section 8.1). Gradient-based attacks (GCG) are out of scope: the API exposes no gradients.

6.3 Seeds and honesty

V, N, and ND re-run $3\times$: standard deviations $\approx$ 0.005 on compliance, 0.000 on ASR and benign success. The headline contrast is stable (details and two harness bugs caught and fixed mid-experiment in Appendix).

6.4 Victim generalization: a 2x2 matrix

All results above use DeepSeek-chat as the victim. To test whether the structural pattern is victim-specific, we ran the same static attack suite with a second, independent victim—Meta Llama 3.1 8B

(local Ollama, same harness, tool registry added to the prompt):

The 2×2 matrix is in Appendix A.2.

Three findings. First, **baseline injection compliance is model-specific**: Llama 8B is more resistant than DeepSeek (compliance 0.328 vs 0.492). Second, N’s **“model luck” is confirmed by a second model**: hardening suppresses actionability to a model-specific degree (0.233 vs 0.014). Third, V’s **structural claim is victim-independent**: 0.000 ASR and 1.000 benign for both, tolerating 0.317 compliance; a willing cross-family attacker (DeepSeek against Llama) also failed (0.000). **Structural guarantees are model-independent; behavioral luck is model-dependent.**

6.5 Validation on the official AgentDojo benchmark

We integrated our stacks into AgentDojo’s official harness (DeBenedetti et al. 2024)—their banking suite, their environment fixtures, four of their attack families (*direct*, *ignore_previous*, *inject_agent*, *system_message*), and their evaluation—with DeepSeek-chat as the agent model. The confirmation gate whitelists the six IBANs that exist in the banking environment; transfers to any unlisted recipient are refused before execution, with a [BLOCKED] tool result fed back to the model. **For a same-setting comparison, we additionally ran AgentDojo’s own defenses (tool_filter, repeat_user_prompt, spotlighting) in the identical harness, model, task subset, and attack set.**

The full benchmark table is in Appendix A.4; the key contrast: V reaches 0.0% ASR at 82.6% utility, while AgentDojo’s tool_filter reaches the same 0.0% at 16.7%.

(95% Wilson CIs in brackets.) Three observations. First, **V matches the strongest official defense on ASR (both 0.0% [0.0, 2.6]) at five times the utility** (82.6% vs 16.7%). Second, **the tool filter’s utility collapse is a VAL failure-mode prediction made visible**: its 16.7% is not “refusing most calls” but DeepSeek failing the filter’s own instruction—it emits tool names not in the suite, so the filter removes every tool (126/144 traces have no tool call). The same defense on GPT-4o *increases* benign utility; on DeepSeek it collapses to 16.7% (to 0% on inspection: the 24 “successes” match a pre-existing transaction). Anchored in the

model’s instruction-following (L0/L1 luck), when the model does not comply the defense destroys usability rather than merely failing to protect. Third, **the cross-testbed difference in N is the strongest evidence for our core claim**: in our testbed N reached 0.000 ASR (model refusal behavior); on AgentDojo the same stack leaks 3.5%. *N’s zero is context-dependent model behavior—move the stack to another benchmark and the zero moves with it.* V’s zero does not move (0.0% in both).

Attack-family breakdown of the undefended baseline. The 6.2% ND ASR is not uniform across the four families: *direct* accounts for 7 of the 9 successes (19.4% ASR) and *system_message* for 2 (5.6%), while *ignore_previous* and *injecagent* reach 0.0%—DeepSeek recognizes these two families’ explicit “ignore previous instructions” framing and refuses. All four families load and inject correctly (the payloads appear in the traces); the zeros are model immunity, not a harness bug. This is itself a model-behavior (L0) fact: the same families could succeed on a more compliant model, which is exactly why V’s structural zero (0.0% across all four) is the meaningful number.

Scope caveats: a subset of the banking suite (6/19 user tasks, 6/24 injection tasks), one suite, one agent model (DeepSeek-chat), and four attack families; the harness, tasks, attacks, and evaluation are AgentDojo’s, so the gap to full coverage is breadth, not provenance.

7 Analysis: The Same Zero, Different Guarantees

The experiment’s central image is two zeros. **V’s 0.000 is structural**: the gate reads platform-recorded events (the LLM cannot write them) and the sandbox enforces schema whitelists; it holds at 0.567 compliance across all 12 attack families, adaptive, and white-box. **N’s 0.000 is behavioral**: the hardening prompt makes this model hedge²⁰¹⁴propose the tool, then refuse the arguments. A different model, stronger attacker, or differently-phrased prompt could move N’s zero; nothing can move V’s within its ODD. The Z(03b1) trajectories make the same point: N’s flatness is the model’s mood holding, V’s is the platform’s record²⁰¹⁴same trajectory, different ground.

This is the deployment answer to “which defense should I buy?”: **VAL selection buys a guarantee**;

intuition buys the model’s current mood. Both reach zero, but N’s zero costs every benign action while V’s is lossless; the AgentDojo tool filter is the same story from the other side^{2014a} defense anchored in the model’s instruction-following *destroys usability* when the model does not comply. Under VAL’s usage criteria, the deployer’s question becomes precise: *is the threat inside the anchor’s ODD?* If yes, L2/L3 structure is available; if no, the honest answer is L2 correctness plus labeling, not a claim of safety.

The account extends to the behavioral layer. **Execution Hallucination (EH)**²⁰¹⁴an agent verbally refuses while the operation completes at the OS level²⁰¹⁴was measured by LITMUS (Zhang et al. 2026) across six agents (EHR 7.98²⁰¹³17.97%), invisible to every semantic-only framework. Under VAL this is a prediction: the anchor for a *behavioral* claim must live on the physical layer; a semantic-layer anchor cannot verify what the model did but did not say (TPR 0.60 (Yin 2026)). ## 8. Limitations

1. **Single model (victim and attacker).** All attacks used DeepSeek-chat; cross-family attackers (Kimi, Llama 8B) declined the role. A less-aligned attacker and GCG remain open variants.
2. **Self-built testbed.** Scenarios, defenses, and sandbox are ours²⁰¹⁴deliberate for a controlled comparison; Section 6.5 validates on AgentDojo’s official harness (0.0% ASR). Our behavioral blind spot (covert execution, TPR 0.60 (Yin 2026)) is independently reproduced by LITMUS (Zhang et al. 2026); we did not run its harness.
3. **Coverage.** Twelve attack families plus adaptive/white-box/PAIR; token-level optimization absent; 20 benign scenarios (CIs in Section 6.1).
4. **Raters and prediction sample.** Blind-validated (security subset 90.9%, 03ba 2248 0.8 (Yin 2026)); 10/10 hits on a small sample²⁰¹⁴the mechanism is the claim, not the hit count.

8 Conclusion

The agent-security field has too many defenses and no pre-purchase axis; we showed that VAL provides one. The anchor predicts how a defense fails²⁰¹⁴and choosing by this axis beat intuition

on the same budget: identical security numbers, opposite guarantees, a 100-point utility gap. We located the L3 frontier in agent security: **confinement, not semantics** encode what restricts what an agent can do, not what judges whether it should. We offer VAL as a falsifiable framework and a research agenda, not a settled ontology.

Appendix

8.1 A.5 Single-defense escalation (static/adaptive/white-box)

8.2 A.4 AgentDojo benchmark results

8.3 A.1 Defense level map

8.4 A.2 Victim generalization (2×2)

8.5 A.3 Varying ND baseline across testbeds

On the varying ND baseline across testbeds.

The undefended ASR differs sharply across our three settings—0.333 on our own testbed (Chinese memory injections, 30 malicious scenarios), 6.2% on AgentDojo (English TODO injections, 144 pairs), 87.5% on JADE’s MCP instances (English tool-description poisoning, 16 cases). This is expected and informative rather than a contradiction: ND ASR measures the *attack surface* (injection strength, carrier, environment semantics), not a fixed property of the model. What is stable across all three is the *contrast*: the intuition stack’s zero moves with the setting (0.000 → 0.035 → 0.062 as attacks get harder), while V’s structural zero does not (0.000 everywhere).

References

1. *CaMeL: Data-Control Separation for LLM Agents (Source-Gated Control Flow)*. 2025. <https://arxiv.org/abs/2503.18813>.
2. Costa, Marcelo et al. 2025. *Securing AI Agents with Information-Flow Control*. <https://arxiv.org/abs/2505.23643>.
3. Debenedetti, Edoardo, Tianjiu Zhang, Mislav Balunovic, Lukas Beurer-Kellner, Marc Fischer, and Florian Tramer. 2024. *Agent-Dojo: A Dynamic Environment to Evaluate Attacks and Defenses for LLM Agents*. <https://arxiv.org/abs/2406.13352>.
4. *Injection-Execution Dissociation: A Mechanistic Evaluation of Persistent Memory Attacks on Stateful LLM Agents*. 2026. <https://arxiv.org/abs/2605.08442>.
5. *LLM-Based Test Oracles: Source-of-Authority Taxonomy—a Systematic Literature Review*. 2026. <https://arxiv.org/abs/2607.05031>.
6. Lu, et al. 2025. *When Does Verification Pay Off? A Closer Look at LLMs as Solution Verifiers*. <https://arxiv.org/abs/2512.02304>.
7. Xu, J., Y. Xiao, W. Shao, H. Liu, and X. Li. 2026. *Memory Provenance Laundering in LLM Agents: A Non-Amplification Firewall for Persistent Memory*. <https://arxiv.org/abs/2607.29167>.
8. Yin, Yajie. 2026. *Grading the Graders: Verification Autonomy Levels (L0–L5) for LLM Reasoning*. <https://arxiv.org/abs/2608.19009>.
9. Zhang, Chiyu, Huiqin Yang, Bendong Jiang, et al. 2026. *LITMUS: Benchmarking Behavioral Jailbreaks of LLM Agents in Real OS Environments*. <https://arxiv.org/abs/2605.10779>.
10. Zhou, Andy et al. 2023. *Defending Against Alignment-Breaking Attacks via Robustly Aligned LLM (RA-LLM)*. <https://arxiv.org/abs/2309.14348>.

Config	Static ASR
D1	0.000
D2	0.200
D3	0.000
D4	0.067
V	0.000
N	0.000

Config	ASR (n = 144,
ND (no defense)	6.2% [3.3, 11.5]
N (prompt hardening + keyword filter)	3.5% [1.5, 7.9]
V (confirmation gate + parameter sandbox)	0.0% [0.0, 2.6]
AgentDojo tool_filter	0.0% [0.0, 2.6]
AgentDojo repeat_user_prompt	8.3% [4.8, 14.0]
AgentDojo spotlighting	6.2% [3.3, 11.5]

Level	Defenses
L0	Prompt hardening
L1	Keyword filter
L2	Llama Guard c
L3	IFC/Fides (Co
L4	Smart-contract
N/A	AgentDojo (De

Config	DeepSeek victi
ND (no defense)	0.333 / 0.492 /
N (D1+D2)	0.000 / 0.233 /
V (D3+D4)	0.000 / 0.497 /